

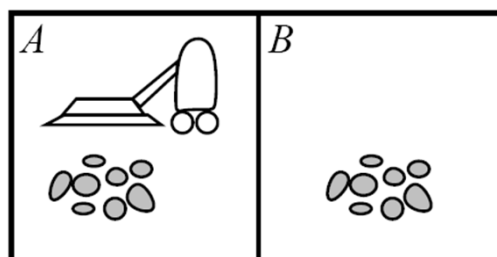
Part 2: Solving Problems by Searching



Univ.-Prof. Dr. Gerhard Widmer
Department of Computational Perception
Johannes Kepler University Linz

gerhard.widmer@jku.at
<http://www.cp.jku.at/people/widmer>

“Motivation”: The Vacuum Cleaner World



“World”: two rooms A, B that may contain dirt

Agent: a vacuum cleaner

Percepts (observations): location and contents (e.g., [A,dirty])

Actions: GoLeft, GoRight, Suck, (NoOp)

Goal: world should be clean

TASK: Write an agent program that, starting in some world configuration S , plans action sequences that lead to achievement of the goal

Overview

Problem solving as search for appropriate action sequences

Definition of search problem

Some simple toy problems

A generic tree search algorithm schema

(Digression: Computational Complexity)

Five uninformed ('blind') search strategies:

- Breadth-First Search
- Uniform-Cost Search
- Depth-First Search (Backtracking Search)
- Depth-limited Search
- Iterative Deepening Search

Solving Problems by Planning Action Sequences

Observations:

- Simple Reflex Agents are too limited to solve real problems:
- Direct mapping from percept sequences to actions is too large to store
- Lack of flexibility and autonomy

More promising: Goal-based Problem Solving Agent

- Wants to achieve a goal
- Can execute a fixed (finite) set of actions
- Knows about the consequences of its actions
- Tries to find sequences of actions that lead to desirable states (goals)
→ no hard-wired program that prescribes every step

“Knowledge Base (KB)”

→ PROBLEM SOLVING as SEARCH for action sequences

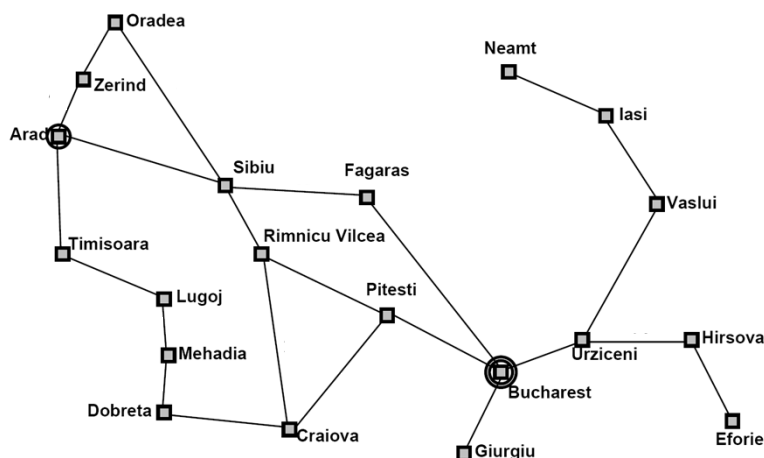
Note: For the moment, assume **“offline”** problem solving: solution is computed once, then executed “with eyes closed”; percepts (other than starting state) are ignored. Assumes that the world is deterministic and actions always produce the desired result

Problem Solving as Search

- The “Problem Solving as Search” paradigm may be one of the most fundamental contributions of (early) AI research to computer science
- AI has developed a large number of different search methods – e.g., “uninformed” search, “heuristic” search, “local” search, stochastic/probabilistic search (e.g., genetic/evolutionary algorithms), ...
- Many of these are now considered standard computer science algorithms
- Many tasks can be formulated as search problems
- Crucial: *appropriate formulation* of the task as a search problem
- Search methods are generic, i.e., independent of a particular task
→ widely and directly applicable

Example: Travelling in Romania – Finding a Way from Arad to Bucharest ¹⁾

Map:



Represented as:

Arad → Zerind
 Zerind → Arad
 Arad → Sibiu
 Arad → Timisoara
 Zerind → Oradea
 Timisoara → Lugoj
 Oradea → Sibiu

¹⁾ from: Russell & Norvig, 2000

Definition: Well-defined Search Problems

A search problem can be formally defined by five components:

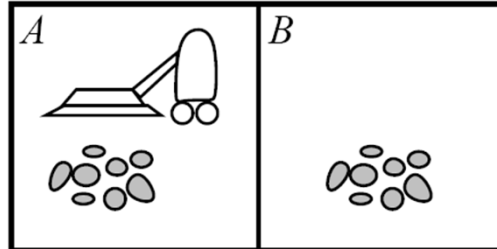
- the complete **set of possible world states** S (e.g., $\{\text{in(Arad)}, \text{in(Zerind)}, \dots\}$)
(terminology: "world states" = whatever is known (through percepts) about the world at a given point in time)
- for each state s , all possible **actions** that can be taken in s , and their **effects** (e.g., in the form of a *successor function* $\text{Succ}(s)$ that gives, for each state s , a set of pairs $\langle \text{action}, \text{resultstate} \rangle$:
 $\text{Succ}(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{in(Zerind)} \rangle, \langle \text{Arad} \rightarrow \text{Sibiu}, \text{in(Sibiu)} \rangle, \dots \}$
- the **starting (initial) state** (e.g., in(Arad))
- a **goal test** (a definition of which states constitute goals); can be
 - explicit (direct description of desired state), e.g., $x = \text{in(Bucharest)}$
 - implicit (needs to be computed), e.g., $\text{checkmate}(x)$
- **costs** $c(x,a)$ of taking action a in state x
(must be non-negative / greater than zero)

Search Problems: Basic Concepts

- **State space** = set of all possible states that can be reached from the initial state (defined by initial state and successor function)
- The state space forms a **graph**
- A **path** in the state space = sequence of states connected by actions
- A **solution** to a problem is a path from the initial state to a goal state
- The **cost** of a solution is the sum of the action costs along the solution path
- An **optimal solution** is a solution with minimum cost
- **Search** = process of finding a (possibly optimal) solution in the state space
- Search generally proceeds by creating (explicitly or implicitly) a **search tree**, exploring a (usually very small) subgraph of the state space graph

Some Simple Example Problems (“Toy Problems”):

1. The Vacuum Cleaner World



“World”: two rooms A, B that may contain dirt

Agent: a vacuum cleaner

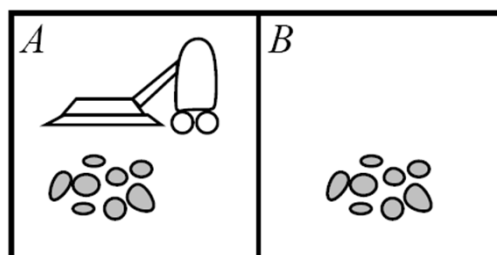
Percepts (observations): location and contents (e.g., [A,dirty])

Actions: GoLeft, GoRight, Suck

Goal: world should be clean

Some Simple Example Problems (“Toy Problems”):

1. The Vacuum Cleaner World



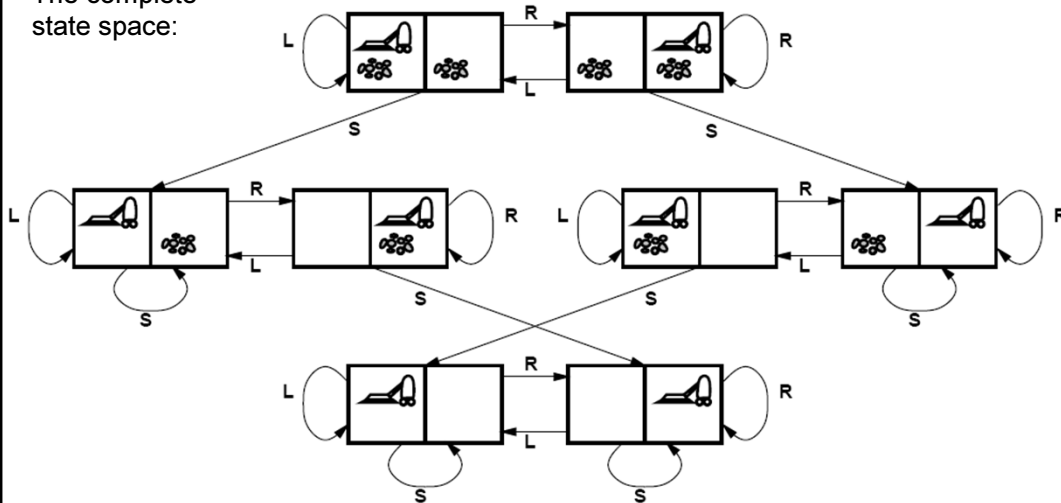
Formalisation as a Search Problem:

- **States:** triplets $\langle \text{RobotLoc}, \text{DirtyOrNotA}, \text{DirtyOrNotB} \rangle$; 2 rooms, each may or may not contain dirt $\rightarrow 2 \times 2 \times 2 = 8$ possible world states
- **Initial state:** any of the states could be the initial state
- **Actions:** Left (L), Right (R), Suck (S)
- **Successor Function:** defines effects of actions (see next slide)
- **Goal Test:** State = $\langle A, \text{clean}, \text{clean} \rangle$ or $\langle B, \text{clean}, \text{clean} \rangle$
- **Action costs:** each step costs 1 \rightarrow optimal solution = shortest path

Some Simple Example Problems (“Toy Problems”):

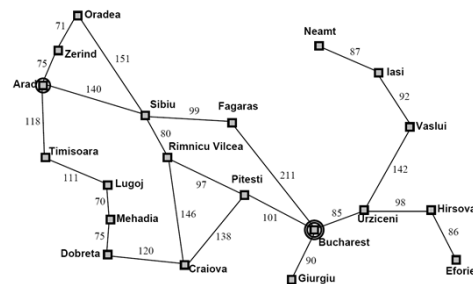
1. The Vacuum Cleaner World

The complete state space:



Some Simple Example Problems (“Toy Problems”):

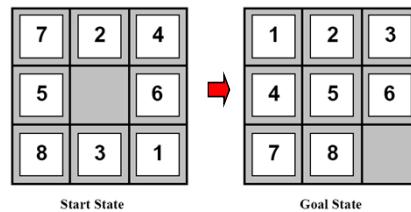
2. The “Traveling in Romania” World



Formalisation as a Search Problem:

- **States:** the agent is in one of N cities $\rightarrow N$ possible world states $\text{in}(x)$
- **Initial state:** $\text{in}(\text{Arad})$
- **Actions:** $x \rightarrow y$, where x is the current city, and y is a city reachable from x via a direct connection
- **Successor Function:** defines all the roads
- **Goal Test:** State = $\text{in}(\text{Bucharest})$
- **Action costs:** length of road from x to y \rightarrow path cost = total trip length

Some Simple Example Problems (“Toy Problems”): 3. The 8-Puzzle



Formalisation as a Search Problem:

- **States:** a state description specifies the location of each of the eight tiles and the blank in one of the 9 squares
 → size of state space: $9!/2 = 181.440$ reachable states!
 (4x4 15-puzzle: 1.300.000.000 ...)
- **Initial state:** any of the states could be the initial state
- **Actions:** MoveBlankLeft (L), MBRight (R), MBUp (U), MBDown(D)
- **Goal Test:** true if numbers obey a certain ordering
- **Action costs:** each step costs 1 → want to find shortest solution

Real-World Examples of Search Problems

- Route Finding
- Robot Navigation and Action Planning
- Automatic Assembly Sequencing, Scheduling
- Protein Design
- Generally: Millions of Optimisation/Configuration Problems
- Internet Searching
- Logical Reasoning (!)
- Learning (!)
- ...

... any problem that can be formalised as finding a sequence of operations to reach a desired state or produce a desired configuration.

Problem Solving as Search

Basic idea of a general search process (simplified):

1. start at initial state
2. "expand" the current state:
try (all) possible actions applicable to current state, producing new states
3. check if one of the new states is a goal
4. if yes → solution found
5. if no → select one of the as yet unexpanded states for new search step
6. go to 2.

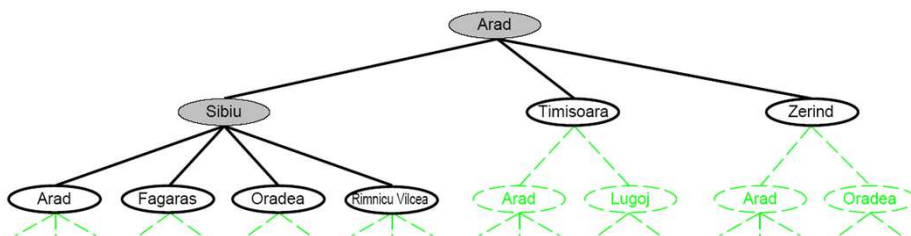
→ Search process produces a **Search Tree**

- **Search Node:** a node in the search tree (i.e., a state with extra information)
- **Expanding** a node/state = generating a new set of states by applying the successor function to the node (i.e., trying out all actions that could be taken in the current state)
- **Search Strategy:** method for deciding which node to expand next

Search Trees

Two more concepts:

- **Leaf Node** = node that has been generated, but not yet expanded
- **„Fringe“** = set of all current leaf nodes



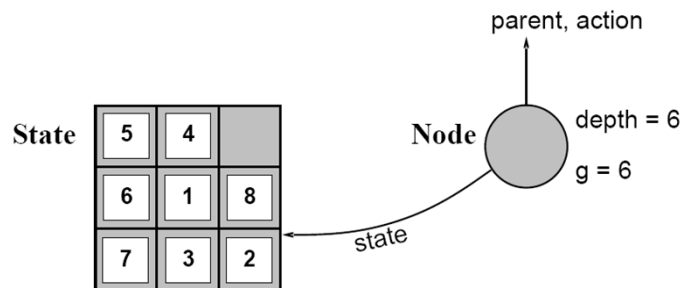
→ **expanded** = {Arad, Sibiu}

→ **fringe** = {Arad, Fagaras, Oradea, Rimnicu, Timisoara, Zerind}

Implementation of Search Nodes

Node = data structure with 5 components:

- **State:** the state in state space to which the node corresponds
- **Parent-Node:** the node in the search tree that generated this node
- **Action:** the action that was applied to the parent to generate the node
- **Path-Cost:** the cost $g(n)$ of the path from the initial state to the node n (the path is given by the parent pointers)
- **Depth:** the number of steps along the path from the initial node



Search Strategy =

Specific way of selecting next node to expand (from fringe)

Implement fringe as a **queue** data structure

Operations on a queue:

- Make-Queue(element,...)
- Empty?(queue)
- First(queue)
- Remove-First(queue)
- Insert(element, queue) ← this is the operation that distinguishes different types of queues ..
- Insert-All(element, ..., queue)

A Generic Tree Search Algorithm

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE(node)) then return node
    fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

Evaluation Criteria for Search Strategies

1. Completeness:

Is the algorithm guaranteed to find a solution if one exists?

2. Optimality:

Is the strategy guaranteed to find the *optimal* solution?

3. Time Complexity:

How long does it take to find a solution?

4. Space Complexity:

How much memory is needed to perform the search?

A Short Digression: Computational Complexity Analysis, $O(\cdot)$ Notation, NP Hard Problems

Complexity Analysis of Algorithms:

Most important questions:

1. How long does algorithm A take to solve problem P ?
(„time complexity“)
2. How much space (memory) does algorithm A require to solve P ?
(„space complexity“)

Complexity Analysis of Algorithms

Two approaches to determining time and space complexity:

1. **Benchmarking** – run the algorithm on problem P (a specific implementation, on a specific machine, given a specific input) and measure the time taken (*ms*) and space used (*bytes*)
→ only very specific information
2. **Asymptotic Analysis** – abstract mathematical analysis of how many „computation steps“ (time) and „storage units“ (space) the algorithm will (roughly) need, given not a specific input, but any input of „size“ n
→ independent of implementation details
→ independent of specific input – tells us how algorithm behaves on a whole *class* of problems (e.g., „sorting lists of length n “)
→ tells us **how computation costs grow with problem size**

Asymptotic Analysis

```

function SUMMATION(sequence) return a number
    sum ← 0
    for i ← 1 to LENGTH(sequence)
        sum ← sum + sequence[i]
    return sum

```

Step 1: abstract over the input:

find a parameter that characterises the size of the input

Example: n = length of the sequence

Step 2: abstract over the algorithm:

find some (implementation-independent) measure that reflects the running time of the algorithm

Example: number of lines of code executed
(or number of additions performed)

Step 3: analyse the algorithm to find a formula T that computes the total number of steps taken by the algorithm as a function of the size of the input

Example: $T(n) = 2n + 2$ (lines of code)

Asymptotic Analysis

Problems with exact analysis:

1. precise number of steps may be different for different inputs
→ usually only possible to compute worst case (T_{worst}) or average case complexity (T_{avg})
2. often impossible to perform exact analysis of the algorithm
→ use „**order approximation**“: „SUMMATION algorithm is $O(n)$ “
Meaning: „its measure $T(n)$ is at most some constant factor times n (with the possible exception of a few small values of n)“

Def.: Order of a function (complexity in „Big O Notation“):

$T(n)$ is $O(f(n))$ if $T(n) \leq k \cdot f(n)$ for some constant k , for all $n > n_0$

Ex.: QUICKSORT is $O(n^2)$ in the worst case, and $O(n \log n)$ in the average case

This is called **asymptotic analysis** (tells us how T grows as n grows)

Complexity Analysis of Problems

Problem complexity analysis:

- analyses the complexity of problems / problem classes (rather than specific algorithms)
- independent of a specific algorithm (i.e., no algorithm can do better in the worst case)

Examples:

- computing a random number between 0 and n is $O(1)$
- finding a word in a dictionary of n words is $O(\log n)$
- computing the maximum in a list of n numbers is $O(n)$
- sorting a list of n numbers is $O(n^2)$
- determining whether two graphs with n vertices are isomorphic is $O(c^n)$ (for some constant c)

Commonly Encountered Orders of Functions / Complexities

notation	name
$O(1)$	constant
$O(\log n)$	logarithmic
$O([\log n]^c)$	polylogarithmic
$O(n)$	linear
$O(n \log n)$	„linearithmic“, quasilinear, supralinear
$O(n^2)$	quadratic
$O(n^c), c > 1$	polynomial, sometimes called „algebraic“
$O(c^n)$	exponential, sometimes called „geometric“
$O(n!)$	factorial

„tractable“

„intractable“

Problem Classes in Algorithmic Complexity Theory

***P* – The class of polynomial problems:**

- solvable in time polynomial in n – $O(n^c)$ for some constant c
- considered „easy“ or „solvable“ („tractable“)
 - contains those problems with running times like $O(n)$ or $O(n \log n)$
 - but also those with complexity $O(n^{1000})$...

***NP* – The class of non-deterministic polynomial problems:**

- definition: a problem is in *NP* if there is some algorithm that can guess a solution and then verify, in polynomial time, whether the solution is correct
- general belief in complexity theory (though still not proved):
 - $NP \neq P$
 - *NP* contains many problems with exponential complexity ($\geq O(c^n)$)
- exponential problems are intractable for even medium-sized inputs

Unfortunately, many of the interesting problems are in *NP* ...

Evaluation Criteria for Search Strategies

1. Completeness:

Is the algorithm guaranteed to find a solution when one exists?

2. Optimality:

Is the strategy guaranteed to find the *optimal* solution?

3. Time Complexity:

How long does it take to find a solution?

4. Space Complexity:

How much memory is needed to perform the search?

➔ *Problem Size* is expressed in terms of 3 quantities:

- b ... branching factor (max. number of successors of any node)
- d ... depth of the shallowest goal node (= length of minimum length solution)
- m ... maximum length of any path in the state space (may be infinite)

Time and space are measured in terms of the number of nodes generated

Five Uninformed (Blind) Search Strategies

- Breadth-First Search
- Uniform-Cost Search
- Depth-First Search (Backtracking Search)
- Depth-limited Search
- Iterative Deepening Search

“uninformed” / “blind”:

- Search strategy has no understanding of the “meaning” of states
- All it can do is generate successors and recognise goal states

Alternative: “informed” / “heuristic” search:

- Search strategy knows (can estimate) which non-goal states are “more promising” than others
- Will be covered in next lecture ...

Five Uninformed (Blind) Search Strategies

... differ in how they instantiate the generic tree search algorithm:

- ➔ which node is expanded next?
- = how is the INSERTALL function realised?

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE(node)) then return node
    fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

Question: Why not test for goal immediately after generating a node?

- ➔ Will understand later: because a better solution might still be found later (before node itself is due for expansion) – cf. uniform cost search etc.; makes no big difference in terms of asymptotic complexity

Breadth-First Search (BFS)

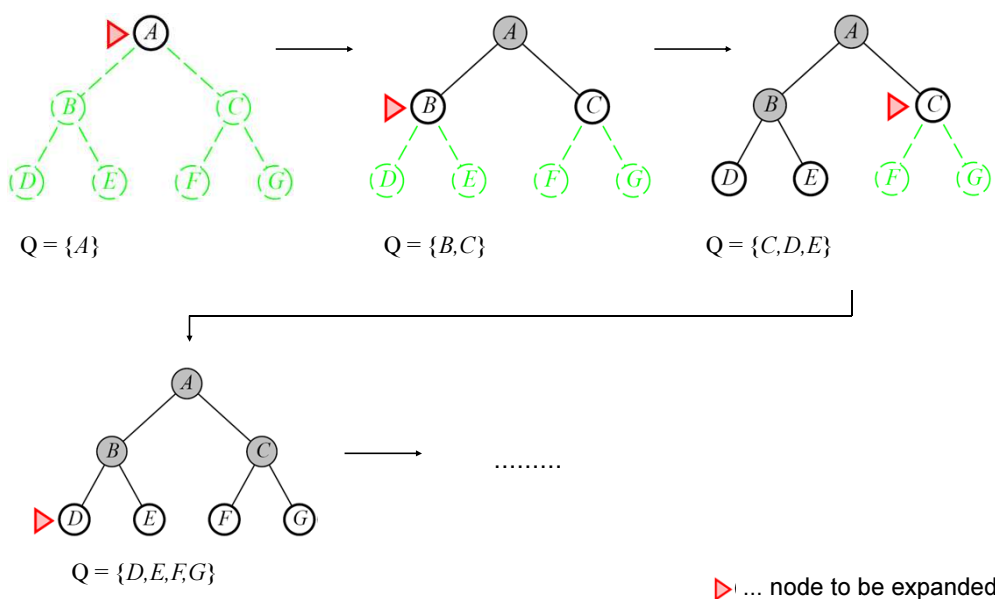
Basic strategy:

- Expand root node (i.e., initial state)
- then expand all successors of the root node
- then expand all their successors, etc.
- In general: expand all nodes at a given depth before expanding any nodes at the next level

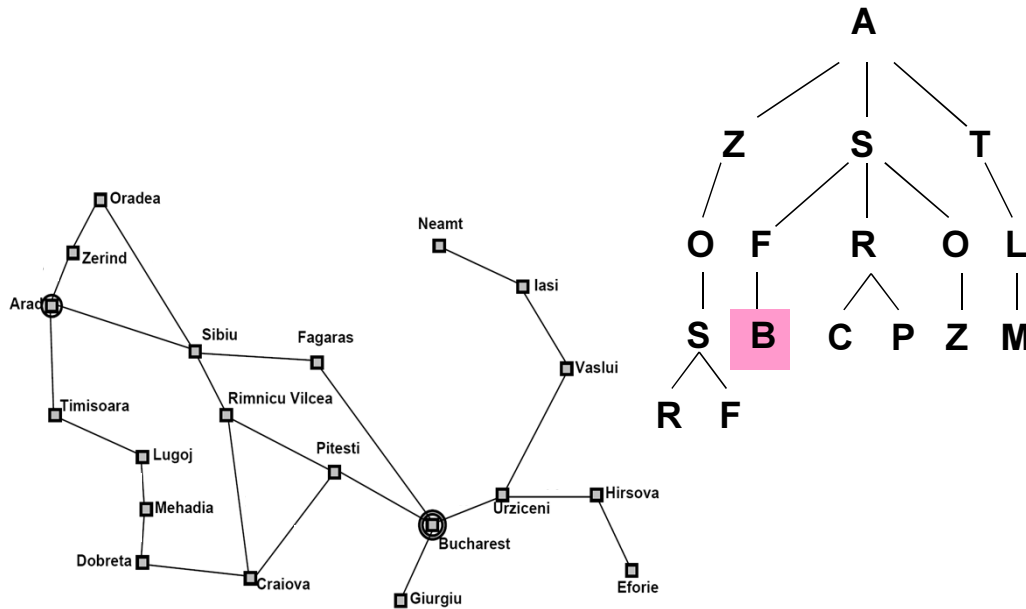
Realisation:

- Call TREE-SEARCH with an empty fringe that is a **first-in-first-out (FIFO) queue**
- FIFO puts all newly generated successors at the end of the queue
 → nodes that are generated first will be expanded first

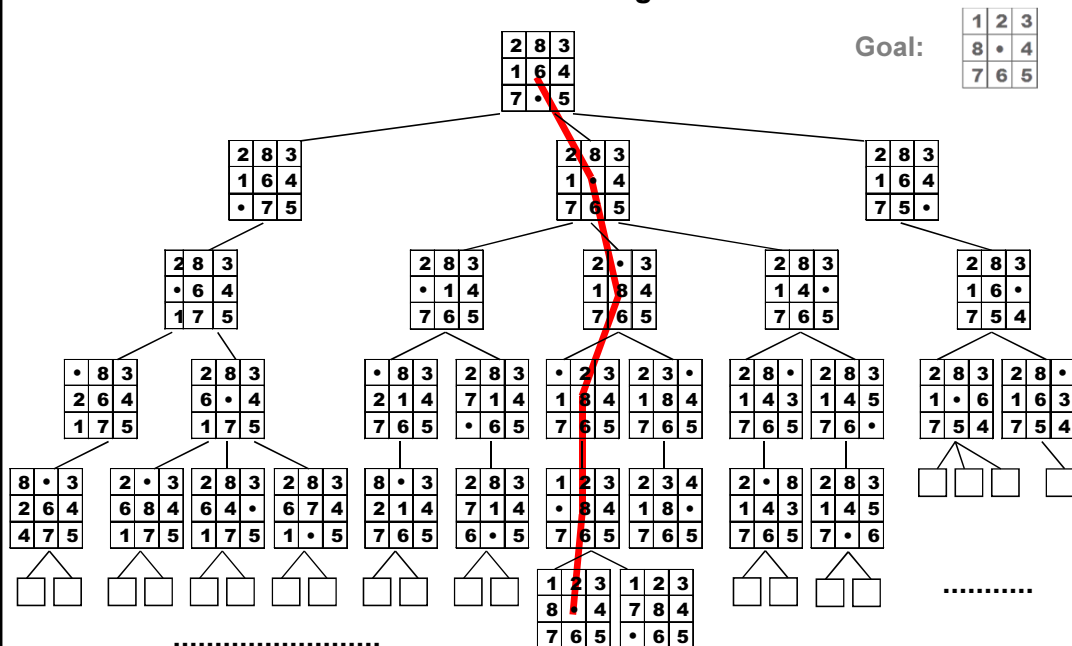
Breadth-First Search (BFS)



Breadth-First Search in Romania



Breadth-First Search: Solving the 8-Puzzle



Properties of Breadth-First Search

Completeness:

BFS is complete – if the shallowest goal node is at depth d , BFS will eventually find it (after expanding all nodes of depth $\leq d$)

Optimality:

BFS is optimal if cost is constant (it first finds the solution with the shortest path); not optimal in general

Time Complexity (in terms of nodes generated):

$$1 + b + b^2 + b^3 + \dots + b^d + (b^{d+1} - b) = O(b^{d+1})$$

Space Complexity:

Fringe must be kept in memory; at depth $d+1$, fringe size is $\leq b^{d+1}$

→ space complexity is $O(b^{d+1})$

b ... branching factor (max. number of successors of any node)
 d ... depth of the shallowest goal node (i.e., goal with minimum length path)
 m ... maximum length of any path in the state space (may be infinite)

Time and Memory Requirements of BFS

Assumptions:

- branching factor $b = 10$
- compute 10,000 nodes/sec
- 1000 bytes / node

$$(\sum_{l=1}^{d+1} b^l) - b$$

Solution depth d	Nodes generated	Time	Memory
2	1,100	0.11 seconds	1 megabyte
4	111,100	11 seconds	100 megabytes
6	10^7	19 minutes	10 gigabytes
8	10^9	31 hours	1 terabyte
10	10^{11}	129 days	100 terabytes
12	10^{13}	35 years	10 petabytes
14	10^{15}	3,523 years	1 exabyte

Breadth-First with Costs: Uniform-Cost Search (UCS)

Generalisation of Breadth-First Search to deal with different action costs

Basic strategy:

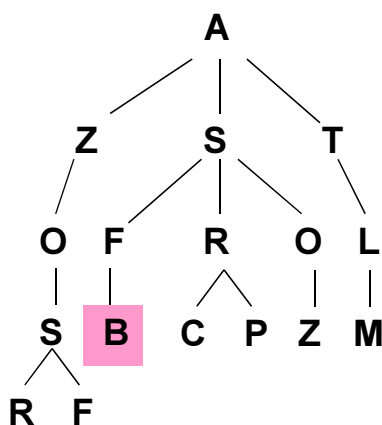
- Expand node with **lowest path cost** so far (i.e., cost from root to node) (instead of node with *shortest* path, as in BFS)
- Effect: finds optimal-cost solution first (instead of shortest-path solution)
- Equivalent to BFS if all costs are equal

Realisation:

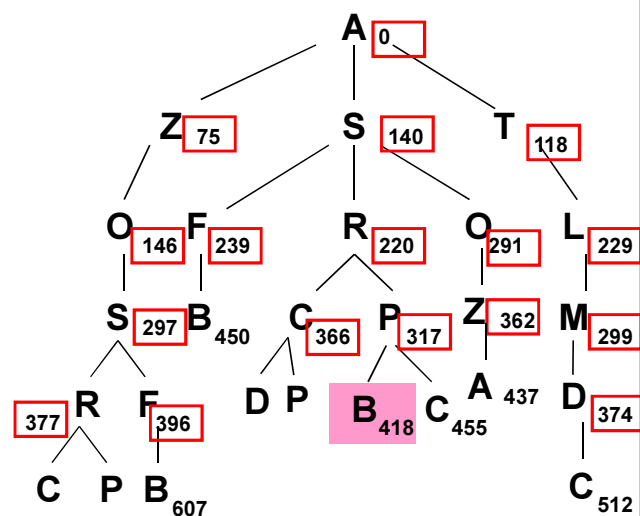
- Fringe = **priority queue** sorted by **path cost** (lowest first)

Uniform-Cost Search in Romania

Breadth-First



Uniform-Cost



Properties of Uniform-Cost Search

Completeness: is complete *if* all actions have cost > 0
(a zero-cost action leading back to the same state (e.g., no-op) would lead to an infinite loop)

Optimality: optimal!

Time and Space Complexity:

- difficult to analyse in terms of d , because search is guided by cost rather than depths
- let C^* = cost of the optimal solution; assume each action costs at least ε
 - worst-case time and space complexity is $O(b^{\lceil C^*/\varepsilon \rceil})$
 - can be much greater than b^d
(because UCS can (and often does) explore large trees of cheap steps before exploring expensive and perhaps useful steps)
 - when all steps have equal costs, then $O(b^{\lceil C^*/\varepsilon \rceil}) = O(b^d)$, of course...

Depth-First Search (DFS)

Basic strategy:

- Always expand the *deepest* node in the current fringe
- Expand a branch until either a solution is found or no more actions possible
- If no more successors (no more actions possible)
 - backtrack to last choice point and expand alternative node

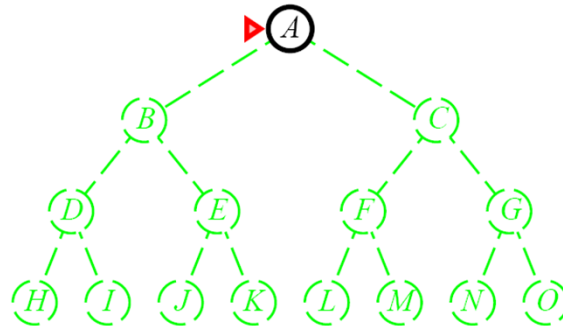
Realisation:

- Fringe = **LIFO (last-in-first-out) queue („stack“)**
(i.e., newly expanded nodes are added to the beginning of the queue)
- Alternative implementation: recursive function that calls itself on each of its children in turn

Advantage:

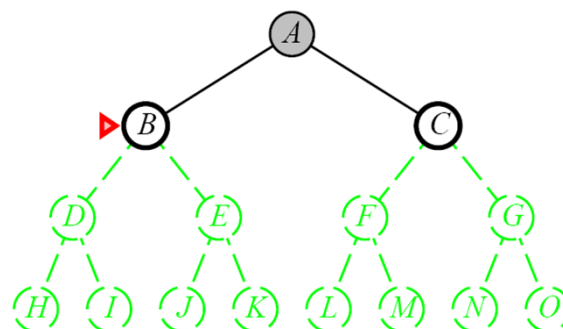
- Low memory requirements: needs to store only a single path at each moment, plus unexpanded siblings for each node on the path
- In other words: the fringe never ‚explodes‘ to exponential size

Depth-First Search



$$Q = \{A\}$$

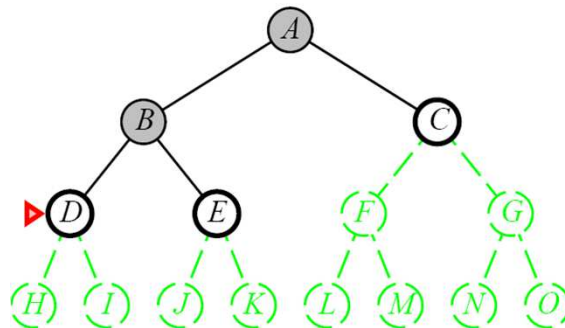
Depth-First Search



$$Q = \{A\}$$

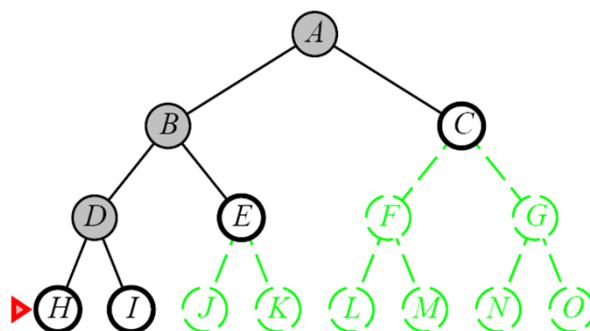
$$Q = \{B, C\}$$

Depth-First Search



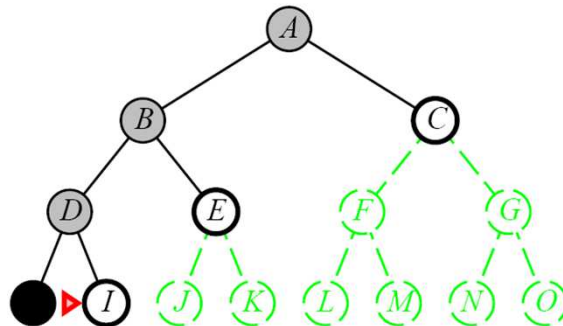
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$

Depth-First Search



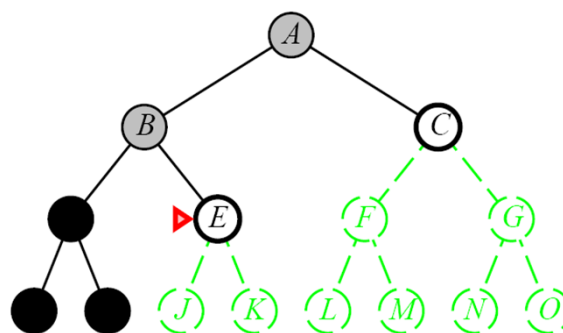
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$

Depth-First Search



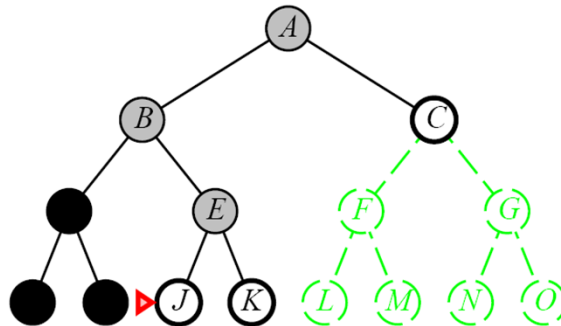
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$

Depth-First Search



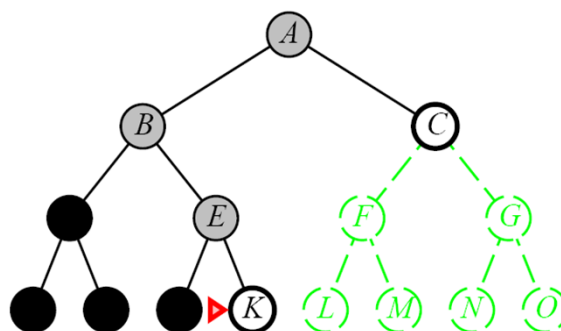
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$
 $Q = \{E, C\}$

Depth-First Search



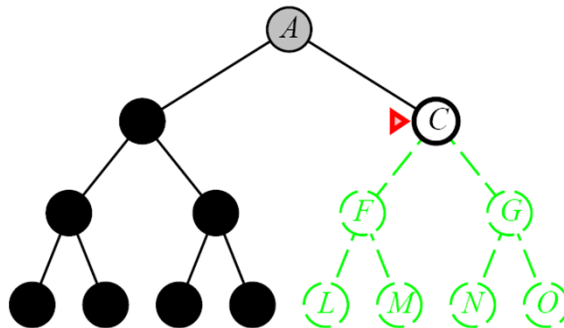
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$
 $Q = \{E, C\}$
 $Q = \{J, K, C\}$

Depth-First Search



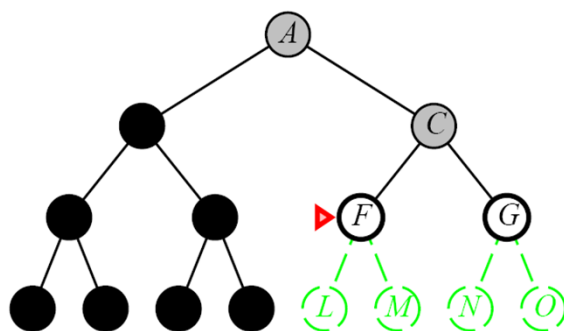
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$
 $Q = \{E, C\}$
 $Q = \{J, K, C\}$
 $Q = \{K, C\}$

Depth-First Search



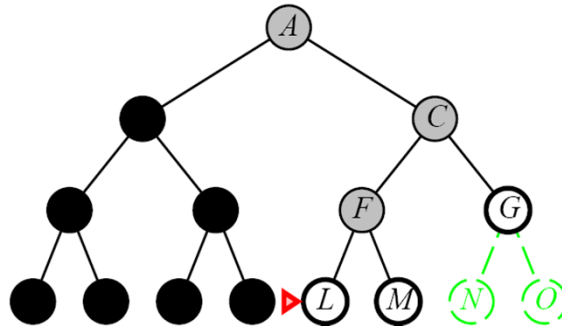
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$
 $Q = \{E, C\}$
 $Q = \{J, K, C\}$
 $Q = \{K, C\}$
 $Q = \{C\}$

Depth-First Search



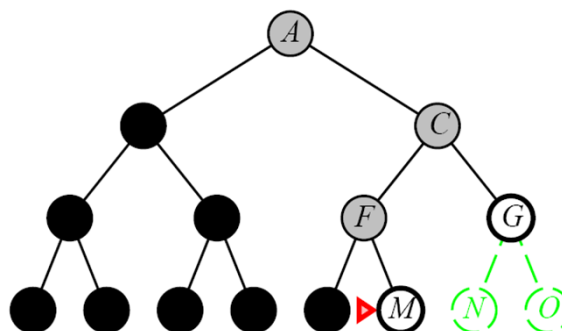
$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$
 $Q = \{E, C\}$
 $Q = \{J, K, C\}$
 $Q = \{K, C\}$
 $Q = \{C\}$
 $Q = \{F, G\}$

Depth-First Search



$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$
 $Q = \{E, C\}$
 $Q = \{J, K, C\}$
 $Q = \{K, C\}$
 $Q = \{C\}$
 $Q = \{F, G\}$
 $Q = \{L, M, G\}$

Depth-First Search



$Q = \{A\}$
 $Q = \{B, C\}$
 $Q = \{D, E, C\}$
 $Q = \{H, I, E, C\}$
 $Q = \{I, E, C\}$
 $Q = \{E, C\}$
 $Q = \{J, K, C\}$
 $Q = \{K, C\}$
 $Q = \{C\}$
 $Q = \{F, G\}$
 $Q = \{L, M, G\}$
 $Q = \{M, G\}$

Properties of Depth-First Search

Completeness:

DFS is **not** complete – can follow infinite paths (if $m = \infty$) and miss solution (even if solution might be very close to root)

Optimality:

DFS is **not** optimal – will find „left-most“ solution, not shortest one

Time Complexity (in terms of nodes generated):

in the worst case, DFS generates entire search tree (if goal is in right-most path)

→ time complexity is $O(b^m) \gg O(b^d) !!$

Space Complexity:

only current branch + unexpanded nodes along branch need to be kept

→ space complexity is $bm + 1 = O(bm)$ → linear space complexity!!

b ... branching factor (max. number of successors of any node)
 d ... depth of the shallowest goal node (i.e., goal with minimum length path)
 m ... maximum length of any path in the state space (may be infinite)

Depth-Limited Search

Goal:

- Fix DFS's problem with unbounded trees (infinite paths)

Basic strategy:

- Use DFS, but with a **fixed depth limit** l
- All nodes at depth l are treated as if they had no successors
- Problem: how to choose an appropriate l ?

Effects:

- l is additional source of incompleteness (if we choose an $l < d$)
- Just like DFS, DLS is non-optimal (even if we choose $l > d$)
- Time complexity is $O(b^l)$
- Space complexity is $O(bl)$

Iterative Deepening Depth-First Search

Basic idea:

- Iterative strategy for finding the best depth limit l
- Repeatedly call Depth-limited DFS with depth limit $l = 0, 1, 2, \dots$, until a goal is found (when depth limit l reaches d)

```

function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution
  inputs: problem, a problem
  for  $depth \leftarrow 0$  to  $\infty$  do
     $result \leftarrow$  DEPTH-LIMITED-SEARCH(problem,  $depth$ )
    if  $result \neq$  cutoff then return  $result$ 
  end
  
```

Iterative Deepening Depth-First Search

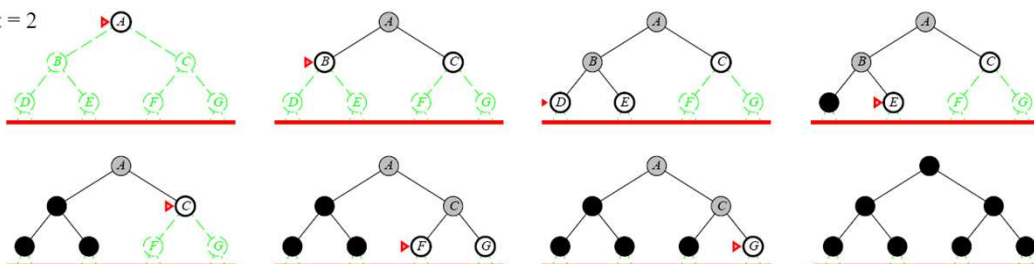
Limit = 0



Limit = 1

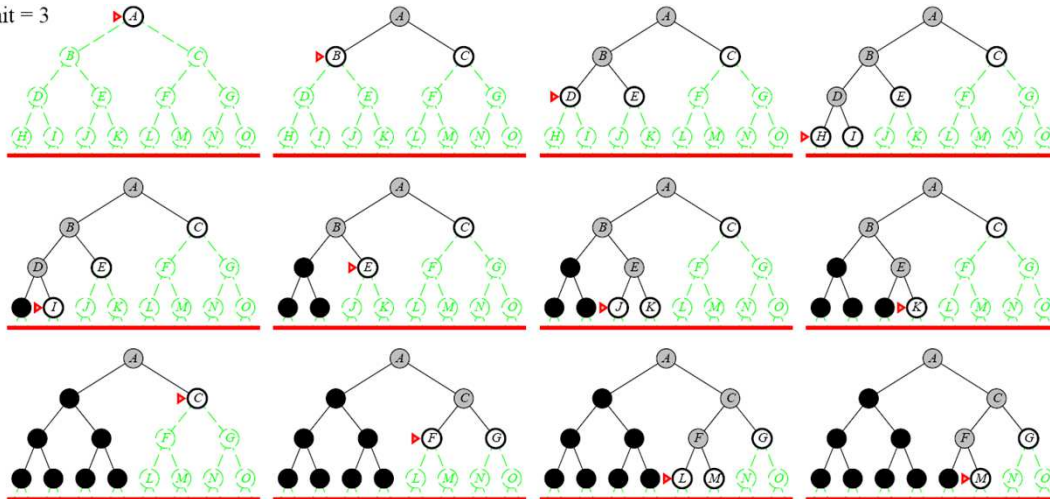


Limit = 2



Iterative Deepening Depth-First Search

Limit = 3



Properties of Iterative Deepening Search

Iterative Deepening combines advantages of DFS and BFS:

Completeness: complete (like breadth-first search)

Optimality: optimal if costs are uniform (like breadth-first search)

Time Complexity (in terms of nodes generated):

$O(b^d)$ (analogous to breadth-first search)

Space Complexity:

$O(bd)$ (analogous to depth-first search)

b ... branching factor (max. number of successors of any node)
 d ... depth of the shallowest goal node (i.e., goal with minimum length path)
 m ... maximum length of any path in the state space (may be infinite)

Isn't Iterative Deepening Wasteful?

Apparent problem: All tree levels $< d$ are generated multiple times!

But:

- Most of the nodes in a search tree (with a fixed branching factor b) are in the bottom level ($\sum_{i=1}^{d-1} b^i < b^d$)
- So it does not matter much if upper levels are generated several times

More precisely:

- In IDS, nodes at bottom level (depth d) are generated 1x, nodes at next-to-bottom level are generated 2x, ..., children of root are generated d x

→ # nodes generated: $(d)b + (d-1)b^2 + \dots + (1)b^d = O(b^d)$

→ Time complexity of IDS is still $O(b^d)$

→ In general, Iterative Deepening is the preferred uninformed search method when there is a large search space and the depth of the solution is unknown

Complexity of Uninformed Search Methods: Summary

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes*	No	Yes, if $l \geq d$	Yes
Time	b^{d+1}	$b^{\lceil C^*/\epsilon \rceil}$	b^m	b^l	b^d
Space	b^{d+1}	$b^{\lceil C^*/\epsilon \rceil}$	bm	bl	bd
Optimal?	Yes*	Yes	No	No	Yes*

b ... branching factor (max. number of successors of any node)
 d ... depth of the shallowest goal node (i.e., goal with minimum length path)
 m ... maximum length of any path in the state space (may be infinite)
 l ... depth limit

Avoiding Repeated States

Problem:

- In many applications, some states can be reached via different action sequences, or actions are reversible (e.g., route finding, 8-puzzle, ...)
- that is, the state space is a *graph*, not a tree
- Consequence: duplication of nodes and subtrees in the search tree, cycles, infinite paths!

Solution:

- Keep track of all nodes already expanded („CLOSED list“)
- Compare newly generated node to nodes in CLOSED list
- Expand new node only if not equal to a node in CLOSED
→ General graph search algorithm

Avoiding Repeated States: The General Graph-Search Algorithm

```

function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
  end
  
```

The General Graph-Search Algorithm

Problems:

- **Optimality** may be lost by simply discarding new, duplicated states (may have lower cost than previously found node with same state)
 - need to compare path costs
 - may need to revise path costs of descendants of kept node
 - things become more complex
- GRAPH-SEARCH must keep every node in memory (in *fringe* or *closed*)
 - depth-first and iterative deepening search no longer have linear space costs
 - back to **$O(b^d)$ space complexity** ...
 - many searches will be impossible because of memory limitations

Problems can sometimes be avoided by clever design of the search problem representation and state space!